AI AUDIT WORKBENCH

Four-Day Hackathon MVP & Submission Plan

Goal
Deliver a small, reliable, end-to-end AI audit workflow with evidence-backed findings, a live deployment, and all required submission materials.

1. Recommended MVP Scope

Build **one complete audit workflow** rather than attempting every global-scale feature. The prototype should demonstrate how a user describes an audit objective, reviews an AI-generated plan, approves execution, and receives findings linked to supporting evidence.

Suggested demo scenario: Illustrative Data Review & Policy Validation

- User enters a generic review objective in a conversational interface.
- Google Gemini generates a structured plan for the requested review.
- User reviews and approves the plan before execution.
- The workflow checks sample records against configurable validation rules or a sample policy document.
- The application flags exceptions and shows the source record, relevant rule, and explanation.
- A final report is generated with findings, evidence references, and execution metadata.

Core experience

Left panel: conversational instructions and clarifications. **Right panel:** workflow stages and live status. **Results:** findings, evidence, and a downloadable report.

Scope guardrails

- Use sample or synthetic data; do not depend on live systems or sensitive production data.
- Keep the workflow small and deterministic where possible; use explicit validation rules for checks.
- Present AI-generated explanations as assistive output and retain human approval for plan execution.
- Describe global regulatory coverage as a future capability unless it is actually implemented in the demo.

2. Four-Day Execution Plan

Prioritize a working vertical slice early. Freeze new features once the end-to-end workflow is stable.

DAY 1 — CORE PIPELINE	Build the audit logic and Gemini integration
Objective	Make the core workflow work from input to findings.
Tasks	Define one generic review scenario; prepare sample records and configurable rules; connect Gemini for plan generation; implement validation checks and structured findings.
Exit criteria	A test input produces a plan, runs the checks, and returns expected findings.
DAY 2 — UI & EVIDENCE	**Make the workflow easy to understand in a demo**
Objective	Show what the system is doing and why a finding was raised.
Tasks	Build chat and workflow status UI; display progress; show each finding with its source and rule; generate a report with timestamps and execution details.
Exit criteria	A teammate can run the demo without developer intervention and understand the results.
DAY 3 — DEPLOY & REPOSITORY	**Deploy a stable public prototype**
Objective	Make the application accessible and reproducible.
Tasks	Deploy to Google Cloud Run or Firebase as appropriate; configure secrets safely; test the public URL; prepare an accessible GitHub repository and README with setup steps.
Exit criteria	The live URL works in a clean browser session and the repository can be accessed by evaluators.
DAY 4 — SUBMISSION & DEMO	**Polish, record, and verify all deliverables**
Objective	Submit a coherent story supported by a working product.
Tasks	Complete the prescribed pitch deck; record a 3–4 minute demo; write the brief solution description; check links, permissions, deployment stability, and submission rules.
Exit criteria	All five deliverables are complete, accessible, and consistent with the actual implementation.

Critical rule: If the core workflow is not stable by the end of Day 2, stop adding features. Use Day 3 to deploy and Day 4 to complete the submission.

3. Global Positioning Without Overbuilding

The global story should be demonstrated through architecture and a small working example, not through a long list of unfinished features.

Capability	What to demonstrate now	Future direction
Adaptable rules	Run the same workflow with a selectable rule set.	Support different jurisdictions and identify control gaps.
Explainable findings	Link findings to source records and the rule that triggered them.	Standardized, independently reviewable evidence packs.
Human oversight	Require approval before the plan executes.	Role-based approvals and enterprise governance.
Reusable workflows	Save and rerun the workflow on another sample dataset.	Reusable industry workflows and a governed agent marketplace.

4. Required Submission Deliverables

1. Pitch deck

Use the prescribed template; explain problem, solution, novelty, architecture, impact, and roadmap.

2. Demo video

3–4 minutes; show the live workflow, a real run, evidence-backed findings, and the result.

3. Live prototype

Deployed link that remains available through evaluation; test it in a clean browser session.

4. GitHub repository

Accessible to evaluators; include README, setup instructions, dependencies, and safe configuration guidance.

5. Brief description

Summarize what was built, who it helps, how AI is used, and what is demonstrated.

Platform requirement: The provided submission instructions require a Google AI model such as Gemini or Gemma, or an agentic platform, and deployment on Google Cloud via Cloud Run or Firebase. Confirm additional rules in the full official guidelines.

5. Suggested 3–4 Minute Demo Storyboard

Time	What to show	Message
0:00–0:25	Problem and generic review objective	Manual review work can be fragmented and slow.
0:25–1:00	Enter objective; Gemini proposes a plan	Natural language becomes an actionable workflow.
1:00–1:25	Review and approve the plan	Human oversight is built in.
1:25–2:25	Run checks; show progress and exceptions	The prototype performs a real workflow.
2:25–3:10	Open a finding and inspect its evidence	Results are traceable and explainable.
3:10–3:40	Show report and global roadmap	The MVP is a foundation for broader assurance.

6. Team Working Principles

- Assign owners for the core workflow, UI, and deployment/submission; combine roles if the team is small.
- Integrate early: agree on API request/response schemas and sample data on Day 1.

- Keep secrets out of GitHub; use environment variables or Google Cloud Secret Manager.
- Use deterministic checks for review decisions; use the LLM for planning and explanation, not as the sole source of truth.
- Rehearse before recording, and keep a backup demo recording in case of network issues.

Success criteria: One generic review objective goes through planning, human approval, execution, and evidence-backed reporting on a stable public deployment — with all five submission deliverables ready.